

# representation\_with\_a\_feature\_cross

August 26, 2021

```
[1]: from jupyterthemes import jtplot
     jtplot.style(theme="onedork", figsize=(14, 5))
```

```
[ ]: #@title Copyright 2020 Google LLC. Double-click for license information.
     # Licensed under the Apache License, Version 2.0 (the "License");
     # you may not use this file except in compliance with the License.
     # You may obtain a copy of the License at
     #
     # https://www.apache.org/licenses/LICENSE-2.0
     #
     # Unless required by applicable law or agreed to in writing, software
     # distributed under the License is distributed on an "AS IS" BASIS,
     # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
     # See the License for the specific language governing permissions and
     # limitations under the License.
```

## 1 Representation with a Feature Cross

In this exercise, you'll experiment with different ways to represent features.

### 1.1 Learning Objectives:

After doing this Colab, you'll know how to:

- Use `tf.feature_column` methods to represent features in different ways.
- Represent features as `bins`.
- Cross bins to create a `feature cross`.

### 1.2 The Dataset

Like several of the previous Colabs, this exercise uses the [California Housing Dataset](#).

### 1.3 Call the import statements

The following code imports the necessary code to run the code in the rest of this Colaboratory.

```
[2]: #@title Load the imports
```

```

# from __future__ import absolute_import, division, print_function,
↳ unicode_literals

import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import feature_column
from tensorflow.keras import layers

from matplotlib import pyplot as plt

# The following lines adjust the granularity of reporting.
pd.options.display.max_rows = 10
pd.options.display.float_format = "{:.1f}".format

tf.keras.backend.set_floatx('float32')

print("Imported the modules.")

```

Imported the modules.

## 1.4 Load, scale, and shuffle the examples

The following code cell loads the separate .csv files and creates the following two pandas DataFrames:

- `train_df`, which contains the training set
- `test_df`, which contains the test set

The code cell then scales the `median_house_value` to a more human-friendly range and then shuffles the examples.

```

[3]: # Load the dataset
train_df = pd.read_csv("https://download.mlcc.google.com/mledu-datasets/
↳ california_housing_train.csv")
test_df = pd.read_csv("https://download.mlcc.google.com/mledu-datasets/
↳ california_housing_test.csv")

# Scale the labels
scale_factor = 1000.0
# Scale the training set's label.
train_df["median_house_value"] /= scale_factor

# Scale the test set's label
test_df["median_house_value"] /= scale_factor

# Shuffle the examples
train_df = train_df.reindex(np.random.permutation(train_df.index))

```

## 1.5 Represent latitude and longitude as floating-point values

Previous Colabs trained on only a single feature or a single synthetic feature. By contrast, this exercise trains on two features. Furthermore, this Colab introduces **feature columns**, which provide a sophisticated way to represent features.

You create feature columns as possible:

- Call a `tf.feature_column` method to represent a single feature, single feature cross, or single synthetic feature in the desired way. For example, to represent a certain feature as floating-point values, call `tf.feature_column.numeric_column`. To represent a certain feature as a series of buckets or bins, call `tf.feature_column.bucketized_column`.
- Assemble the created representations into a Python list.

A neighborhood's location is typically the most important feature in determining a house's value. The California Housing dataset provides two features, `latitude` and `longitude` that identify each neighborhood's location.

The following code cell calls `tf.feature_column.numeric_column` twice, first to represent `latitude` as floating-point value and a second time to represent `longitude` as floating-point values.

This code cell specifies the features that you'll ultimately train the model on and how each of those features will be represented. The transformations (collected in `fp_feature_layer`) don't actually get applied until you pass a DataFrame to it, which will happen when we train the model.

```
[4]: # Create an empty list that will eventually hold all feature columns.
feature_columns = []

# Create a numerical feature column to represent latitude.
latitude = tf.feature_column.numeric_column("latitude")
feature_columns.append(latitude)

# Create a numerical feature column to represent longitude.
longitude = tf.feature_column.numeric_column("longitude")
feature_columns.append(longitude)

# Convert the list of feature columns into a layer that will ultimately become
# part of the model. Understanding layers is not important right now.
fp_feature_layer = layers.DenseFeatures(feature_columns)
```

When used, the layer processes the raw inputs, according to the transformations described by the feature columns, and packs the result into a numeric array. (The model will train on this numeric array.)

## 1.6 Define functions that create and train a model, and a plotting function

The following code defines three functions:

- `create_model`, which tells TensorFlow to build a linear regression model and to use the `feature_layer_as_fp` as the representation of the model's features.
- `train_model`, which will ultimately train the model from training set examples.
- `plot_the_loss_curve`, which generates a loss curve.

```
[5]: #@title Define functions to create and train a model, and a plotting function
def create_model(my_learning_rate, feature_layer):
    """Create and compile a simple linear regression model."""
    # Most simple tf.keras models are sequential.
    model = tf.keras.models.Sequential()

    # Add the layer containing the feature columns to the model.
    model.add(feature_layer)

    # Add one linear layer to the model to yield a simple linear regressor.
    model.add(tf.keras.layers.Dense(units=1, input_shape=(1,)))

    # Construct the layers into a model that TensorFlow can execute.
    model.compile(optimizer=tf.keras.optimizers.
→RMSprop(learning_rate=my_learning_rate),
                loss="mean_squared_error",
                metrics=[tf.keras.metrics.RootMeanSquaredError()])

    return model

def train_model(model, dataset, epochs, batch_size, label_name):
    """Feed a dataset into the model in order to train it."""

    features = {name:np.array(value) for name, value in dataset.items()}
    label = np.array(features.pop(label_name))
    history = model.fit(x=features, y=label, batch_size=batch_size,
                        epochs=epochs, shuffle=True)

    # The list of epochs is stored separately from the rest of history.
    epochs = history.epoch

    # Isolate the mean absolute error for each epoch.
    hist = pd.DataFrame(history.history)
    rmse = hist["root_mean_squared_error"]

    return epochs, rmse

def plot_the_loss_curve(epochs, rmse):
    """Plot a curve of loss vs. epoch."""

    plt.figure()
    plt.xlabel("Epoch")
    plt.ylabel("Root Mean Squared Error")

    plt.plot(epochs, rmse, label="Loss")
```

```
plt.legend()
plt.ylim([rmse.min()*0.94, rmse.max()* 1.05])
plt.show()

print("Defined the create_model, train_model, and plot_the_loss_curve functions.
↪")
```

Defined the create\_model, train\_model, and plot\_the\_loss\_curve functions.

## 1.7 Train the model with floating-point representations

The following code cell calls the functions you just created to train, plot, and evaluate a model.

```
[6]: # The following variables are the hyperparameters.
learning_rate = 0.05
epochs = 30
batch_size = 100
label_name = 'median_house_value'

# Create and compile the model's topography.
my_model = create_model(learning_rate, fp_feature_layer)

# Train the model on the training set.
epochs, rmse = train_model(my_model, train_df, epochs, batch_size, label_name)

plot_the_loss_curve(epochs, rmse)

print("\n: Evaluate the new model against the test set:")
test_features = {name:np.array(value) for name, value in test_df.items()}
test_label = np.array(test_features.pop(label_name))
my_model.evaluate(x=test_features, y=test_label, batch_size=batch_size)
```

Epoch 1/30

```
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
```

Consider rewriting this model with the Functional API.

```
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
```

```
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
```

Consider rewriting this model with the Functional API.

```
170/170 [=====] - 4s 1ms/step - loss: 48967.7785 -
root_mean_squared_error: 213.1454
```

Epoch 2/30

```
170/170 [=====] - 0s 1ms/step - loss: 13577.8688 -
root_mean_squared_error: 116.5125
```

Epoch 3/30

```
170/170 [=====] - 0s 1ms/step - loss: 13212.2934 -
root_mean_squared_error: 114.9330
```

Epoch 4/30

```
170/170 [=====] - 0s 1ms/step - loss: 13494.9081 -
root_mean_squared_error: 116.1651
```

Epoch 5/30

```
170/170 [=====] - 0s 1ms/step - loss: 13662.1923 -
root_mean_squared_error: 116.8796
```

Epoch 6/30

```
170/170 [=====] - 0s 1ms/step - loss: 13080.7060 -
root_mean_squared_error: 114.3563
```

Epoch 7/30

```
170/170 [=====] - 0s 1ms/step - loss: 13298.3145 -
root_mean_squared_error: 115.3080
```

Epoch 8/30

```
170/170 [=====] - 0s 1ms/step - loss: 13312.5290 -
root_mean_squared_error: 115.3703
```

Epoch 9/30

```
170/170 [=====] - 0s 1ms/step - loss: 13093.2639 -
root_mean_squared_error: 114.4116
```

Epoch 10/30

```
170/170 [=====] - 0s 1ms/step - loss: 13153.6313 -
root_mean_squared_error: 114.6867
```

Epoch 11/30

```
170/170 [=====] - 0s 1ms/step - loss: 12949.4399 -
root_mean_squared_error: 113.7914
```

Epoch 12/30

```
170/170 [=====] - 0s 1ms/step - loss: 13460.4567 -
root_mean_squared_error: 116.0098
```

Epoch 13/30

```
170/170 [=====] - 0s 1ms/step - loss: 13309.8032 -
root_mean_squared_error: 115.3636
```

Epoch 14/30

```
170/170 [=====] - 0s 1ms/step - loss: 13071.6555 -
root_mean_squared_error: 114.3288
```

Epoch 15/30  
170/170 [=====] - 0s 1ms/step - loss: 12903.1461 -  
root\_mean\_squared\_error: 113.5825

Epoch 16/30  
170/170 [=====] - 0s 1ms/step - loss: 12975.5904 -  
root\_mean\_squared\_error: 113.9045

Epoch 17/30  
170/170 [=====] - 0s 1ms/step - loss: 13140.3450 -  
root\_mean\_squared\_error: 114.6255

Epoch 18/30  
170/170 [=====] - 0s 1ms/step - loss: 13404.0652 -  
root\_mean\_squared\_error: 115.7643

Epoch 19/30  
170/170 [=====] - 0s 1ms/step - loss: 13043.3169 -  
root\_mean\_squared\_error: 114.2046

Epoch 20/30  
170/170 [=====] - 0s 1ms/step - loss: 12829.7842 -  
root\_mean\_squared\_error: 113.2589

Epoch 21/30  
170/170 [=====] - 0s 1ms/step - loss: 12946.8356 -  
root\_mean\_squared\_error: 113.7780

Epoch 22/30  
170/170 [=====] - 0s 1ms/step - loss: 12973.4305 -  
root\_mean\_squared\_error: 113.8938

Epoch 23/30  
170/170 [=====] - 0s 1ms/step - loss: 13020.4727 -  
root\_mean\_squared\_error: 114.1027

Epoch 24/30  
170/170 [=====] - 0s 1ms/step - loss: 12923.4633 -  
root\_mean\_squared\_error: 113.6781

Epoch 25/30  
170/170 [=====] - 0s 1ms/step - loss: 13132.5606 -  
root\_mean\_squared\_error: 114.5908

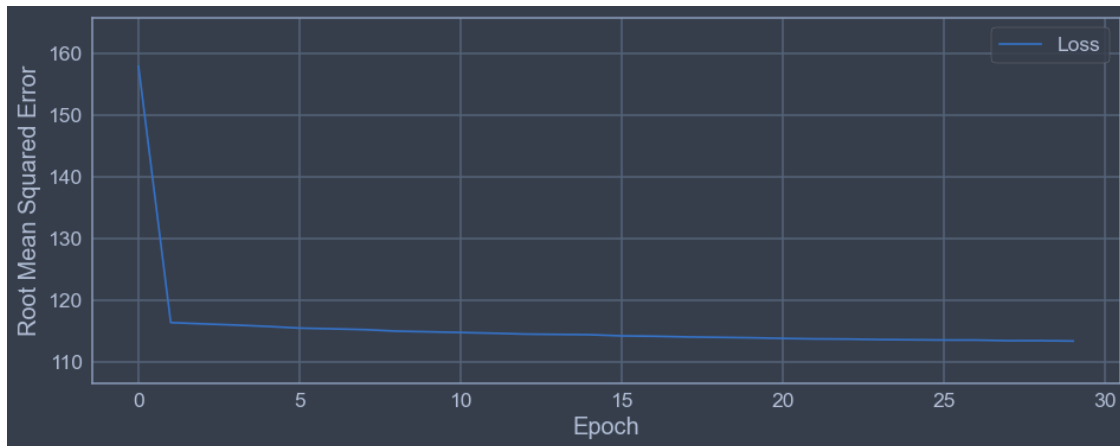
Epoch 26/30  
170/170 [=====] - 0s 1ms/step - loss: 12766.6098 -  
root\_mean\_squared\_error: 112.9832

Epoch 27/30  
170/170 [=====] - 0s 1ms/step - loss: 12964.8454 -  
root\_mean\_squared\_error: 113.8580

Epoch 28/30  
170/170 [=====] - 0s 1ms/step - loss: 12933.9153 -  
root\_mean\_squared\_error: 113.7139

Epoch 29/30  
170/170 [=====] - 0s 1ms/step - loss: 12627.9517 -  
root\_mean\_squared\_error: 112.3555

Epoch 30/30  
170/170 [=====] - 0s 1ms/step - loss: 13061.5998 -  
root\_mean\_squared\_error: 114.2813



: Evaluate the new model against the test set:

```
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
```

Consider rewriting this model with the Functional API.

```
30/30 [=====] - 0s 1ms/step - loss: 12190.5361 -
root_mean_squared_error: 110.4108
```

[6]: [12190.5361328125, 110.41075897216797]

## 1.8 Task 1: Why aren't floating-point values a good way to represent latitude and longitude?

Are floating-point values a good way to represent latitude and longitude?

[ ]: *#@title Double-click to view an answer to Task 1.*

```
# No. Representing latitude and longitude as
# floating-point values does not have much
# predictive power. For example, neighborhoods at
# latitude 35 are not 36/35 more valuable
# (or 35/36 less valuable) than houses at
# latitude 36.
```



```
# Representing `latitude` and `longitude` as  
# floating-point values provides almost no  
# predictive power. We're only using the raw values  
# to establish a baseline for future experiments  
# with better representations.
```

## 1.9 Represent latitude and longitude in buckets

The following code cell represents latitude and longitude in buckets (bins). Each bin represents all the neighborhoods within a single degree. For example, neighborhoods at latitude 35.4 and 35.8 are in the same bucket, but neighborhoods in latitude 35.4 and 36.2 are in different buckets.

The model will learn a separate weight for each bucket. For example, the model will learn one weight for all the neighborhoods in the “35” bin, a different weight for neighborhoods in the “36” bin, and so on. This representation will create approximately 20 buckets:

- 10 buckets for latitude.
- 10 buckets for longitude.

```
[7]: resolution_in_degrees = 1.0  
  
# Create a new empty list that will eventually hold the generated feature_  
↪column.  
feature_columns = []  
  
# Create a bucket feature column for latitude.  
latitude_as_a_numeric_column = tf.feature_column.numeric_column("latitude")  
latitude_boundaries = list(np.arange(int(min(train_df['latitude'])),  
                                   int(max(train_df['latitude'])),  
                                   resolution_in_degrees))  
latitude = tf.feature_column.bucketized_column(latitude_as_a_numeric_column,  
                                                latitude_boundaries)  
feature_columns.append(latitude)  
  
# Create a bucket feature column for longitude.  
longitude_as_a_numeric_column = tf.feature_column.numeric_column("longitude")  
longitude_boundaries = list(np.arange(int(min(train_df['longitude'])),  
                                   int(max(train_df['longitude'])),  
                                   resolution_in_degrees))  
longitude = tf.feature_column.bucketized_column(longitude_as_a_numeric_column,  
                                                longitude_boundaries)  
feature_columns.append(longitude)  
  
# Convert the list of feature columns into a layer that will ultimately become  
# part of the model. Understanding layers is not important right now.  
buckets_feature_layer = layers.DenseFeatures(feature_columns)
```

## 1.10 Train the model with bucket representations

Run the following code cell to train the model with bucket representations rather than floating-point representations:

```
[8]: # The following variables are the hyperparameters.
learning_rate = 0.04
epochs = 35

# Build the model, this time passing in the buckets_feature_layer.
my_model = create_model(learning_rate, buckets_feature_layer)

# Train the model on the training set.
epochs, rmse = train_model(my_model, train_df, epochs, batch_size, label_name)

plot_the_loss_curve(epochs, rmse)

print("\n: Evaluate the new model against the test set:")
my_model.evaluate(x=test_features, y=test_label, batch_size=batch_size)
```

Epoch 1/35

```
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
```

Consider rewriting this model with the Functional API.

```
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
```

Consider rewriting this model with the Functional API.

```
170/170 [=====] - 1s 2ms/step - loss: 53864.0429 -
root_mean_squared_error: 232.0660
```

Epoch 2/35

```
170/170 [=====] - 0s 1ms/step - loss: 46282.1519 -
root_mean_squared_error: 215.1198
```

Epoch 3/35  
170/170 [=====] - 0s 1ms/step - loss: 39309.7831 -  
root\_mean\_squared\_error: 198.2604

Epoch 4/35  
170/170 [=====] - 0s 1ms/step - loss: 33877.7666 -  
root\_mean\_squared\_error: 184.0291

Epoch 5/35  
170/170 [=====] - 0s 1ms/step - loss: 28549.9910 -  
root\_mean\_squared\_error: 168.9535

Epoch 6/35  
170/170 [=====] - 0s 2ms/step - loss: 23718.3898 -  
root\_mean\_squared\_error: 153.9988

Epoch 7/35  
170/170 [=====] - 0s 1ms/step - loss: 20639.1274 -  
root\_mean\_squared\_error: 143.6427

Epoch 8/35  
170/170 [=====] - 0s 2ms/step - loss: 17035.4874 -  
root\_mean\_squared\_error: 130.5014

Epoch 9/35  
170/170 [=====] - 0s 1ms/step - loss: 14806.4360 -  
root\_mean\_squared\_error: 121.6740

Epoch 10/35  
170/170 [=====] - 0s 2ms/step - loss: 13232.4208 -  
root\_mean\_squared\_error: 115.0056

Epoch 11/35  
170/170 [=====] - 0s 2ms/step - loss: 12307.3044 -  
root\_mean\_squared\_error: 110.9284

Epoch 12/35  
170/170 [=====] - 0s 2ms/step - loss: 11496.0458 -  
root\_mean\_squared\_error: 107.2111

Epoch 13/35  
170/170 [=====] - 0s 2ms/step - loss: 10967.1314 -  
root\_mean\_squared\_error: 104.7201

Epoch 14/35  
170/170 [=====] - 0s 1ms/step - loss: 10896.3460 -  
root\_mean\_squared\_error: 104.3580

Epoch 15/35  
170/170 [=====] - 0s 1ms/step - loss: 10675.8825 -  
root\_mean\_squared\_error: 103.3191


Epoch 16/35  
170/170 [=====] - 0s 2ms/step - loss: 10454.8923 -  
root\_mean\_squared\_error: 102.2429

Epoch 17/35  
170/170 [=====] - 0s 1ms/step - loss: 10633.7093 -  
root\_mean\_squared\_error: 103.1159

Epoch 18/35  
170/170 [=====] - 0s 1ms/step - loss: 10577.0247 -  
root\_mean\_squared\_error: 102.8419

Epoch 19/35  
170/170 [=====] - 0s 2ms/step - loss: 10478.8993 -  
root\_mean\_squared\_error: 102.3543  
Epoch 20/35  
170/170 [=====] - 0s 1ms/step - loss: 10594.9817 -  
root\_mean\_squared\_error: 102.9278  
Epoch 21/35  
170/170 [=====] - 0s 1ms/step - loss: 10650.3214 -  
root\_mean\_squared\_error: 103.1961  
Epoch 22/35  
170/170 [=====] - 0s 1ms/step - loss: 10623.2918 -  
root\_mean\_squared\_error: 103.0575  
Epoch 23/35  
170/170 [=====] - 0s 1ms/step - loss: 10443.9517 -  
root\_mean\_squared\_error: 102.1854  
Epoch 24/35  
170/170 [=====] - 0s 1ms/step - loss: 10314.2343 -  
root\_mean\_squared\_error: 101.5552  
Epoch 25/35  
170/170 [=====] - 0s 1ms/step - loss: 10348.0475 -  
root\_mean\_squared\_error: 101.7225  
Epoch 26/35  
170/170 [=====] - 0s 1ms/step - loss: 10183.0090 -  
root\_mean\_squared\_error: 100.9049  
Epoch 27/35  
170/170 [=====] - 0s 1ms/step - loss: 10454.0507 -  
root\_mean\_squared\_error: 102.2394  
Epoch 28/35  
170/170 [=====] - 0s 1ms/step - loss: 10255.0272 -  
root\_mean\_squared\_error: 101.2531  
Epoch 29/35  
170/170 [=====] - 0s 1ms/step - loss: 10355.7453 -  
root\_mean\_squared\_error: 101.7452  
Epoch 30/35  
170/170 [=====] - 0s 1ms/step - loss: 10585.1057 -  
root\_mean\_squared\_error: 102.8796  
Epoch 31/35  
170/170 [=====] - 0s 2ms/step - loss: 10429.9680 -  
root\_mean\_squared\_error: 102.1255  
Epoch 32/35  
170/170 [=====] - 0s 3ms/step - loss: 10358.2837 -  
root\_mean\_squared\_error: 101.7723  
Epoch 33/35  
170/170 [=====] - 0s 3ms/step - loss: 10294.1062 -  
root\_mean\_squared\_error: 101.4425  
Epoch 34/35  
170/170 [=====] - 0s 3ms/step - loss: 10359.1232 -  
root\_mean\_squared\_error: 101.7778

```
Epoch 35/35
170/170 [=====] - 0s 2ms/step - loss: 10424.9198 -
root_mean_squared_error: 102.0992
```



```
: Evaluate the new model against the test set:
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
Consider rewriting this model with the Functional API.
30/30 [=====] - 0s 2ms/step - loss: 10198.6270 -
root_mean_squared_error: 100.9883
```

```
[8]: [10198.626953125, 100.98825073242188]
```

## 1.11 Task 2: Did buckets outperform floating-point representations?

Compare the model's `root_mean_squared_error` values for the two representations (floating-point vs. buckets)? Which model produced lower losses?

```
[ ]: #@title Double-click for an answer to Task 2.
```

```
# Bucket representation outperformed
# floating-point representations.
# However, you can still do far better.
```

### 1.12 Task 3: What is a better way to represent location?

Buckets are a big improvement over floating-point values. Can you identify an even better way to identify location with `latitude` and `longitude`?

```
[ ]: #@title Double-click to view an answer to Task 3.

# Representing location as a feature cross should
# produce better results.

# In Task 2, you represented latitude in
# one-dimensional buckets and longitude in
# another series of one-dimensional buckets.
# Real-world locations, however, exist in
# two dimension. Therefore, you should
# represent location as a two-dimensional feature
# cross. That is, you'll cross the 10 or so latitude
# buckets with the 10 or so longitude buckets to
# create a grid of 100 cells.

# The model will learn separate weights for each
# of the cells.
```

### 1.13 Represent location as a feature cross

The following code cell represents location as a feature cross. That is, the following code cell first creates buckets and then calls `tf.feature_column.crossed_column` to cross the buckets.

```
[17]: resolution_in_degrees = 0.4

# Create a new empty list that will eventually hold the generated feature_
# column.
feature_columns = []

# Create a bucket feature column for latitude.
latitude_as_a_numeric_column = tf.feature_column.numeric_column("latitude")
latitude_boundaries = list(np.arange(int(min(train_df['latitude'])),
    int(max(train_df['latitude'])), resolution_in_degrees))
latitude = tf.feature_column.bucketized_column(latitude_as_a_numeric_column,
    latitude_boundaries)

# Create a bucket feature column for longitude.
longitude_as_a_numeric_column = tf.feature_column.numeric_column("longitude")
longitude_boundaries = list(np.arange(int(min(train_df['longitude'])),
    int(max(train_df['longitude'])), resolution_in_degrees))
longitude = tf.feature_column.bucketized_column(longitude_as_a_numeric_column,
    longitude_boundaries)
```

```

# Create a feature cross of latitude and longitude.
latitude_x_longitude = tf.feature_column.crossed_column([latitude, longitude],
↳hash_bucket_size=100)
crossed_feature = tf.feature_column.indicator_column(latitude_x_longitude)
feature_columns.append(crossed_feature)

# Convert the list of feature columns into a layer that will later be fed into
# the model.
feature_cross_feature_layer = layers.DenseFeatures(feature_columns)

```

Invoke the following code cell to test your solution for Task 2. Please ignore the warning messages.

```

[18]: # The following variables are the hyperparameters.
learning_rate = 0.04
epochs = 35

# Build the model, this time passing in the feature_cross_feature_layer:
my_model = create_model(learning_rate, feature_cross_feature_layer)

# Train the model on the training set.
epochs, rmse = train_model(my_model, train_df, epochs, batch_size, label_name)

plot_the_loss_curve(epochs, rmse)

print("\n: Evaluate the new model against the test set:")
my_model.evaluate(x=test_features, y=test_label, batch_size=batch_size)

```

Epoch 1/35

WARNING:tensorflow:Layers in a Sequential model should only have a single input tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor 'ExpandDims\_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor 'ExpandDims\_2:0' shape=(100, 1) dtype=float32>, 'housing\_median\_age': <tf.Tensor 'ExpandDims\_1:0' shape=(100, 1) dtype=float32>, 'total\_rooms': <tf.Tensor 'ExpandDims\_7:0' shape=(100, 1) dtype=float32>, 'total\_bedrooms': <tf.Tensor 'ExpandDims\_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor 'ExpandDims\_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor 'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median\_income': <tf.Tensor 'ExpandDims\_4:0' shape=(100, 1) dtype=float32>}

Consider rewriting this model with the Functional API.

WARNING:tensorflow:Layers in a Sequential model should only have a single input tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor 'ExpandDims\_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor 'ExpandDims\_2:0' shape=(100, 1) dtype=float32>, 'housing\_median\_age': <tf.Tensor 'ExpandDims\_1:0' shape=(100, 1) dtype=float32>, 'total\_rooms': <tf.Tensor 'ExpandDims\_7:0' shape=(100, 1) dtype=float32>, 'total\_bedrooms': <tf.Tensor 'ExpandDims\_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor 'ExpandDims\_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor 'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median\_income': <tf.Tensor

```

'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
Consider rewriting this model with the Functional API.
170/170 [=====] - 1s 2ms/step - loss: 54619.3328 -
root_mean_squared_error: 233.6996
Epoch 2/35
170/170 [=====] - 0s 2ms/step - loss: 50226.7875 -
root_mean_squared_error: 224.1012
Epoch 3/35
170/170 [=====] - 0s 2ms/step - loss: 45037.2616 -
root_mean_squared_error: 212.2039
Epoch 4/35
170/170 [=====] - 0s 2ms/step - loss: 42016.5040 -
root_mean_squared_error: 204.9477
Epoch 5/35
170/170 [=====] - 0s 2ms/step - loss: 37625.7285 -
root_mean_squared_error: 193.9441
Epoch 6/35
170/170 [=====] - 0s 2ms/step - loss: 33634.8849 -
root_mean_squared_error: 183.3722
Epoch 7/35
170/170 [=====] - 0s 2ms/step - loss: 30226.6555 -
root_mean_squared_error: 173.8476
Epoch 8/35
170/170 [=====] - 0s 3ms/step - loss: 26821.2219 -
root_mean_squared_error: 163.7662
Epoch 9/35
170/170 [=====] - 1s 3ms/step - loss: 24001.8744 -
root_mean_squared_error: 154.9179
Epoch 10/35
170/170 [=====] - 1s 4ms/step - loss: 21735.1485 -
root_mean_squared_error: 147.4095
Epoch 11/35
170/170 [=====] - 1s 4ms/step - loss: 19754.9901 -
root_mean_squared_error: 140.5310
Epoch 12/35
170/170 [=====] - 1s 4ms/step - loss: 18155.7249 -
root_mean_squared_error: 134.7181
Epoch 13/35
170/170 [=====] - 1s 3ms/step - loss: 16080.9936 -
root_mean_squared_error: 126.7975
Epoch 14/35
170/170 [=====] - 1s 4ms/step - loss: 14836.0994 -
root_mean_squared_error: 121.7662
Epoch 15/35
170/170 [=====] - 1s 3ms/step - loss: 13467.5284 -
root_mean_squared_error: 116.0459
Epoch 16/35
170/170 [=====] - 1s 3ms/step - loss: 12560.5400 -

```




root\_mean\_squared\_error: 112.0666  
Epoch 17/35  
170/170 [=====] - 1s 3ms/step - loss: 11591.2872 -  
root\_mean\_squared\_error: 107.6579  
Epoch 18/35  
170/170 [=====] - 1s 3ms/step - loss: 10974.3131 -  
root\_mean\_squared\_error: 104.7483  
Epoch 19/35  
170/170 [=====] - 1s 3ms/step - loss: 10384.7929 -  
root\_mean\_squared\_error: 101.8976  
Epoch 20/35  
170/170 [=====] - 1s 3ms/step - loss: 10099.9753 -  
root\_mean\_squared\_error: 100.4937  
Epoch 21/35  
170/170 [=====] - 1s 3ms/step - loss: 10086.0510 -  
root\_mean\_squared\_error: 100.4264  
Epoch 22/35  
170/170 [=====] - 1s 3ms/step - loss: 9563.2488 -  
root\_mean\_squared\_error: 97.7832  
Epoch 23/35  
170/170 [=====] - 1s 3ms/step - loss: 9769.0690 -  
root\_mean\_squared\_error: 98.8362  
Epoch 24/35  
170/170 [=====] - 1s 3ms/step - loss: 9549.9982 -  
root\_mean\_squared\_error: 97.7176  
Epoch 25/35  
170/170 [=====] - 1s 3ms/step - loss: 9178.3903 -  
root\_mean\_squared\_error: 95.7912  
Epoch 26/35  
170/170 [=====] - 1s 3ms/step - loss: 9493.9538 -  
root\_mean\_squared\_error: 97.4233  
Epoch 27/35  
170/170 [=====] - 1s 3ms/step - loss: 9254.3647 -  
root\_mean\_squared\_error: 96.1959  
Epoch 28/35  
170/170 [=====] - 1s 3ms/step - loss: 9069.2396 -  
root\_mean\_squared\_error: 95.2273  
Epoch 29/35  
170/170 [=====] - 1s 3ms/step - loss: 9031.9349 -  
root\_mean\_squared\_error: 95.0347  
Epoch 30/35  
170/170 [=====] - 1s 3ms/step - loss: 8775.9093 -  
root\_mean\_squared\_error: 93.6712  
Epoch 31/35  
170/170 [=====] - 1s 3ms/step - loss: 8892.9110 -  
root\_mean\_squared\_error: 94.2939  
Epoch 32/35  
170/170 [=====] - 1s 3ms/step - loss: 8922.0123 -

```

root_mean_squared_error: 94.4510
Epoch 33/35
170/170 [=====] - 1s 3ms/step - loss: 8809.8588 -
root_mean_squared_error: 93.8578
Epoch 34/35
170/170 [=====] - 1s 3ms/step - loss: 8873.4916 -
root_mean_squared_error: 94.1924
Epoch 35/35
170/170 [=====] - 1s 3ms/step - loss: 8679.1867 -
root_mean_squared_error: 93.1578

```



```

: Evaluate the new model against the test set:
WARNING:tensorflow:Layers in a Sequential model should only have a single input
tensor, but we receive a <class 'dict'> input: {'longitude': <tf.Tensor
'ExpandDims_3:0' shape=(100, 1) dtype=float32>, 'latitude': <tf.Tensor
'ExpandDims_2:0' shape=(100, 1) dtype=float32>, 'housing_median_age': <tf.Tensor
'ExpandDims_1:0' shape=(100, 1) dtype=float32>, 'total_rooms': <tf.Tensor
'ExpandDims_7:0' shape=(100, 1) dtype=float32>, 'total_bedrooms': <tf.Tensor
'ExpandDims_6:0' shape=(100, 1) dtype=float32>, 'population': <tf.Tensor
'ExpandDims_5:0' shape=(100, 1) dtype=float32>, 'households': <tf.Tensor
'ExpandDims:0' shape=(100, 1) dtype=float32>, 'median_income': <tf.Tensor
'ExpandDims_4:0' shape=(100, 1) dtype=float32>}
Consider rewriting this model with the Functional API.
30/30 [=====] - 0s 3ms/step - loss: 8645.8350 -
root_mean_squared_error: 92.9830

```

[18]: [8645.8349609375, 92.98297882080078]

## 1.14 Task 4: Did the feature cross outperform buckets?

Compare the model's `root_mean_squared_error` values for the two representations (buckets vs. feature cross)? Which model produced lower losses?

```
[ ]: #@title Double-click for an answer to this question.
```

```
# Yes, representing these features as a feature  
# cross produced much lower loss values than  
# representing these features as buckets
```

### 1.15 Task 5: Adjust the resolution of the feature cross

Return to the code cell in the “Represent location as a feature cross” section. Notice that `resolution_in_degrees` is set to 1.0. Therefore, each cell represents an area of 1.0 degree of latitude by 1.0 degree of longitude, which corresponds to a cell of 110 km by 90 km. This resolution defines a rather large neighborhood.

Experiment with `resolution_in_degrees` to answer the following questions:

1. What value of `resolution_in_degrees` produces the best results (lowest loss value)?
2. Why does loss increase when the value of `resolution_in_degrees` drops below a certain value?

Finally, answer the following question:

3. What feature (that does not exist in the California Housing Dataset) would be a better proxy for location than latitude X longitude.

```
[ ]: #@title Double-click for possible answers to Task 5.
```

```
#1. A resolution of ~0.4 degree provides the best  
# results.
```

```
#2. Below ~0.4 degree, loss increases because the  
# dataset does not contain enough examples in  
# each cell to accurately predict prices for  
# those cells.
```

```
#3. Postal code would be a far better feature  
# than latitude X longitude, assuming that  
# the dataset contained sufficient examples  
# in each postal code.
```